

Batch Active Learning for Text Classification and Sentiment Analysis

Abstract

The performance of text classification with supervised models is tied to quality and diversity of the data. The process of data collection and labeling may involve a lot of resources. The intuitive and the most standard approach is to sequentially extend a dataset with labeled data until reaching satisfactory metrics. Active learning techniques optimize the process of sequential unlabeled data selection, so that the annotations would provide the most information about the dataset. The problem of active learning becomes more complex when the sampling is done in batches. In this paper we show a study of advanced batch sampling techniques on text data and the problem of text classification and sentiment analysis. The study compares i) baseline algorithm based on agglomerative instance clustering with the subsequent sampling from clusters given minimum margin of class probabilities ii) warm start modifications of baseline techniques, and iii) Bayesian active learning baseline modification thanks to their ability of better representation of the classification uncertainty. The latter method in warm-start version too.

Transformers encoders show the state-of-the-art results in majority of NLP tasks. In this article, we use RoBERTa for text encoding. The methods are tested on three types datasets, context integrity (Kaggle Gibberish dataset), fake news detection (Kaggle Fake News Detection dataset) and sentiment classification (Twitter Sentiment140 and Amazon Review Classification datasets). We show that both warm-start and Bayesian baseline algorithm modifications outperform the state-of-art approach.

1 Introduction

The development of a text classifier on a new problem requires the availability of the training data and their labels. Labeling involves human annotators and a common practice is to label as many text documents

as possible, train a classifier and search for new data and labels if the performance is unsatisfactory. Random choice of the documents for the data set extension can be costly because the new documents may not bring new information for the classification. Active learning strategy aims to select among available unlabeled documents those that the classifier is most uncertain about and queries an annotator for their labels. Therefore, it has the potential to greatly reduce the effort needed for the development of a new system. While it was introduced almost two decades ago, recent improvements in deep learning motivate our attempt to revisit the topic. For example, SVM-based active learning approaches for text classification date back to 2001 [22], where the superiority of active learning over random sampling was demonstrated. Since deep recurrent and convolutional neural networks achieve better classification results, Bayesian active learning methods for deep networks gained popularity especially in image classification [10], [15]. The Bayesian approach is concerned with querying labels for data for which the classifier predicts the greatest uncertainty. The uncertainty is quantified using the so-called acquisition function, such as predictive entropy [18], mutual information, etc.. Despite the strengths of uncertainty based models, clustering or gradient prediction based algorithms also reach state-of-art results for different tasks.

While different acquisition functions often provide similar results, different representations of predictive distribution yield much more diverse results. The most popular approach using Dropout MC [10] has been tested on text classification [1] and named entity recognition [19], [15], however other techniques such as Langevin dynamics [26] and deep ensembles [13] are available. Deep ensembles often achieve better performance [3], [21] but require higher computational cost since they train an ensemble of networks after each extension of the data set. One potential solution of this problem has been recently proposed in [23], where the ensemble is not trained from a fresh random initialization after each query but initialized randomly around the position of the ensembles from the previous iteration. In this contribution, we test agglomerative clustering state-of-art approach with Bayesian networks, warm-start modifications and different acquisition functions. Our research is underlaid with dif-

ferent types of text data, classification tasks complexity analysis and a sensitivity study for the batch size choice.

Active learning for fake news detection was considered in [4] using uncertainty based on probability of classification. It was later extended to a context aware approach [5]. An entropy based approach has been presented in [12] using an ensemble of three different kinds of networks.

2 Methods

Throughout the paper, we will use RoBERTa base [14] embedding algorithms. RoBERTa is a modified BERT transformer model [6] based on multi-head attention layers [24]. Transformer models provide state of the art results both in context understanding and active learning text classification problems. Representation of the i -th text document \mathbf{x}_i is calculated as the mean value from sentence embeddings of all sentences in the text

$$\mathbf{x}_i = \frac{1}{|D_i|} \sum_{j \in D_i} f_{\text{Sentence embed}}(C^{(j)})$$

where D_i is the set of vectors where each vector represents a sentence in the i -th document, $|D_i|$ is a cardinality of D_i , $C^{(j)}$ is a matrix of j -th sentence where words are encoded with byte pair encoding [20] technique and $f_{\text{Sentence embed}}$ is a function that creates sentence embeddings with respect to the given byte pair encoded words. Sentence embeddings for RoBERTa are output of deep neural network models.

For supervised classification, each document embedding \mathbf{x}_i has to have an associated label \mathbf{y}_i . We are concerned with binary classification for simplicity, however, an extension to multiclass is straightforward. We assume that for the full corpus of text documents $\mathbf{X} = [\mathbf{x}_1, \dots, \mathbf{x}_n]$, only an initial set of $l_0 \ll n$ labels is available, $\mathbf{y}^{(0)} = [y_{l_1}, \dots, y_{l_0}]$, splitting the full set \mathbf{X} to the labeled, $\mathbf{X}^{(0)} = [\mathbf{x}_{l_1}, \dots, \mathbf{x}_{l_0}]$, and unlabeled parts, $\mathbf{X} \setminus \mathbf{X}^{(0)}$. Active learning is defined as a sequential extension of the training data set. In each iteration, $l = 1, \dots, L$, the algorithm computes the value of acquisition function for batch of documents in the unlabeled dataset and selects the indices of documents given some criteria (results of acquisition function), formally:

$$\mathbf{k}_l = A(\mathbf{y}_{k_1}(\theta), \dots, \mathbf{y}_{k_b}(\theta), \mathbf{x}_{k_1}, \dots, \mathbf{x}_{k_b}),$$

where $\mathbf{k}_l = \{k_{l_1}, \dots, k_{l_b}\} \subset \mathbf{K}$ is a subset of batch size b of indices of all unlabeled documents \mathbf{K} , $\mathbf{y}(\theta) = \{y_{\theta_1}, \dots, y_{\theta_d}\}$ are the predictions given classifier parameters empirical distribution of size d trained on all labeled data $p(\theta | \mathbf{X}^{(l-1)}, \mathbf{y}^{(l-1)})$. The documents of the selected indices are sent to the human annotator with a request for labeling. When the selected texts are annotated, the texts are added with its label to the la-

beled data set $\mathbf{X}^l = [\mathbf{X}^{(l-1)}, \mathbf{x}_{k_1}^{(l)}, \dots, \mathbf{x}_{k_b}^{(l)}]$, $\mathbf{y}^{(l)} = [\mathbf{y}^{(l-1)}, y_{k_1}^{(l)}, \dots, y_{k_b}^{(l)}]$.

In this article we do a detailed comparison between a combination of active learning algorithms and acquisition functions.

2.1 Active Learning Methods

The key component of the method is a representation of the posterior distribution of the parameter θ . Due to the complexity of the neural networks it is always represented by samples, with a different method of their generation. We will compare the following methods: i) Dropout MC: is an extension of the ordinary dropout that samples binary mask multiplying output of a layer, hence stopping propagation through all neurons where zeros is sampled through the network. The extension applies the sampled mask even for predictions generating samples from the predictive distribution [10], ii) Deep ensembles: consist of d networks trained in parallel from different initial conditions [13], and iii) Softmax uncertainty: is the most simple approach that uses only one network to estimate a single value θ [4]. Ensembles based approach is the current state-of-the-art in active learning [3]. While many of these have been tested in active learning, the majority of authors assumed that after each step of active learning, the network training starts from the initial conditions. This is clearly suboptimal, since the information from the previous training is lost. A simple solution was presented in [23], where it was argued that estimated results from the previous step can be used as centroids around which the new initial point is sampled. Since this is a form of a warm-start, we also test batch active learning warm-start strategies for Dropout and ensembles. In the authors present a broad analysis for one-instance sampling active learning and show that previously chosen techniques work the best. The methods for representation of parametric uncertainty are:

2.1.1 Deep Ensemble Filter (DEnFi)

is a deep ensemble method with 5 neural networks in the ensembles and warm-start training strategy [13] using weights of the ensemble members in the previous iteration as initial conditions for the new ensemble. Each weight is perturbed by an additive Gaussian noise of variance $q = 0.3$ which is a hyperparameter and given it is the optimal choice that performs good in both short and long active learning runs. In our experiments, the ensemble is trained with parameters `initialization_epochs = 2500` on the initial data and with additional `warm_start_epochs = 700` epochs after each extension of the learning data set.

2.1.2 Dropout MC

is the standard algorithm [10] that trains only a single network with sampled dropout indices and uses the sampling even in the prediction step. Generation of the Monte Carlo prediction is obtained by sampling different values of the dropout binary variable and one for-

ward pass of the network for each sample. We study warm-start with weights from the previous iteration perturbed by an additive noise of variance $q = 0.3$ with 700 epochs. Dropout rate is 0.2. The noise variance choice is driven with same as in the section above.

2.1.3 Softmax Uncertainty

The simplest approach to uncertainty representation is a single neural network with a softmax output layer that considers uncertainty as the output of the softmax score. The method has been applied to active learning in [4] and given the cooperative analysis from performed the worst from the previous models and cold-start algorithms that are trained from scratch after each active learning iteration. The softmax uncertainty is the type of superstructure used on the top of encoding model in state-of-art followed with an acquisition function described in further sections.

2.2 Acquisition Functions

The text classification problem is formed for binary categories. Hence, all generalizations are provided to binary classification tasks.

2.2.1 HAC Entropy

The method is based on clustering the instances with hierarchical agglomerative clustering (HAC) with average linkage, and choosing the smallest non-singleton clusters with the minimum prediction margin values. The sampling algorithm in a combination with softmax uncertainty method performs state-of-art active learning results in . The formal representations is:

$$\tilde{k}_{l_i} = \operatorname{argmax}_{\tilde{k} \in \mathbf{K}} \mathbb{E}_{p(\theta|\mathbf{X}^{(l-1)}, \mathbf{Y}^{(l-1)})} (\mathbb{H}(y_{\tilde{k}}|\theta)), i \in \{l_1, \dots, l_M\} \quad (1)$$

where $\mathbb{H}(y_k|\theta)$ is conditional entropy of the predicted class for \mathbf{x}_k . The batch that will be used for sampling from clustered data is created as:

$$\tilde{\mathbf{k}}_l = \{\tilde{k}_{l_1}, \dots, \tilde{k}_{l_M}\},$$

are $M > b$ values with the highest entropy sampled independently from the unlabeled set. Next, we apply HAC method to create $\mathcal{C} = \{\mathbf{C}_1, \dots, \mathbf{C}_N\}$ clusters and find $m \leq M$ the smallest non-singleton clusters with values $\mathbf{X}_{\tilde{\mathbf{k}}_l}$. Furthermore, $\mathbf{X}_{\mathbf{k}_l}$, $\mathbf{k}_l = \{k_{l_1}, \dots, k_{l_b}\}$ is sampled from \mathcal{C}_m . The sampling is provided from different clusters. However, if $m < b$ the values can be samples from the same cluster.

2.2.2 HAC Min Margin

The HAC Min Margin method was shown in and is the same as shown in 1 where θ is a point-wise estimate. Hence the expectation step \mathbb{E} can be removed from the equation. The next steps of the action sequence for \mathbf{k}_l creation are the same.

2.3 HAC BALD

HAC BALD method repeats the sampling part from 2.2.1 section. The difference in $\tilde{\mathbf{k}}_l$ step. The authors of

the BALD method propose mutual information maximization as follows:

$$\begin{aligned} \tilde{k}_{l_i} &= \operatorname{argmax}_{\tilde{k} \in \mathbf{K}} \mathbb{I}(y_{\tilde{k}}|\mathbf{X}^{(l-1)}, \mathbf{Y}^{(l-1)}), i \in \{l_1, \dots, l_M\} \\ &= \operatorname{argmax}_{\tilde{k} \in \mathbf{K}} \left(\mathbb{H}(y_{\tilde{k}}|\mathbf{X}^{(l-1)}, \mathbf{Y}^{(l-1)}) - \mathbb{E}_{p(\theta|\mathbf{X}^{(l-1)}, \mathbf{Y}^{(l-1)})} (\mathbb{H}(y_{\tilde{k}}|\theta)) \right), \end{aligned} \quad (2)$$

where

$$\mathbb{H}(y_{\tilde{k}}|\mathbf{X}^{(l-1)}, \mathbf{Y}^{(l-1)}) = \mathbb{H}(\mathbb{E}_{p(\theta|\mathbf{X}^{(l-1)}, \mathbf{Y}^{(l-1)})} y_{\tilde{k}}(\theta)).$$

2.4 Entropy

Entropy sampling is generalized as $\mathbf{k}_l = \tilde{\mathbf{k}}_l$, where $\tilde{\mathbf{k}}_l$ is calculated from 1 with $M = b$.

2.5 BALD

Similarly as we derived entropy sampling in section 2.4, we will derive BALD sampling as $\mathbf{k}_l = \tilde{\mathbf{k}}_l$, where $\tilde{\mathbf{k}}_l$ is calculated from 2 with $M = b$.

Described approach was published half a decade ago, and it was modified for batch sampling in and . The method of calculating pseudo mutual information in is slow. The authors of HAC Min Margin method also referred to its inefficiency due to too complex computations. Second approach from that does not have the blocker in computation time and is fast as Entropy and Bald sampling did not show good results during our methods evaluation and bigger computation rounds preparation. Hence, the original BALD method is the only one that we will present in results.

3 Experiments

3.1 Simulation

The methods were compared on different datasets and different batch size. The used datasets are positive/negative tweets from the Tweets [11], Fake News Detection [8], two pairs of categories from Amazon Reviews and Gibberish datasets. The batch sizes are 10, 20, 50 and 100 instances per active learning or random sampling iteration. Specifically, we compared active learning and random sampling strategies for different settings of algorithms, batch sizes and different representations of uncertainty. Each experiment was initiated by random choice of the initial training set of $l_0 = 10$ samples from 10000 text documents (5000 text documents per category), which were the initial 10000 documents of the datasets given categories. For each experiment we continue the active learning simulation until we sample 1000 with querying iterations. Hence, L ranges from 10 to 100 requests for annotation. The batch selection follows the ϵ -greedy approach [25], i.e. the samples given the acquisition function is accepted with probability $\epsilon = \frac{\exp(l-3)}{\exp(l-3)+1}$. A batch of random documents is selected for labeling if not accepted. After each request, the classification performance is evaluated on the remaining part of the selected dataset (i.e.

on the 9990 text documents in the first evaluation) using the area under the ROC curve (AUC) metrics [9]. In order to make the results statistically valid, we repeat the described simulation loop 5 times for all datasets.

3.2 Methods

We compare results for 12 different tuples of algorithms and acquisition functions that include i) five MC Dropout simulations with all acquisition functions, random sampling, all datasets, and all batch sizes, ii) three NN Warm-Start runs based on dropout algorithm with point-wise parameters distribution estimate for HAC Entropy, Entropy, random sampling acquisition functions, all datasets, and all batch sizes, iii) five DEnFi simulations with all acquisition functions, random sampling, three datasets, and only one batch size (due to the computational complexity), and iv) HAC Min-margin run with softmax uncertainty cold-start active learning strategy and HAC Entropy computed for all datasets and all batch sizes. The cold start strategy choice for HAC Min-margin is based on the experimental results from [1], where the hot start-methods (fine tuning without noise perturbation) scored the worst. Hence, for the sake of not having computational biases we decided to train HAC Min-margin algorithm from scratch in every active learning iteration.

3.3 Text Classification

The comparison of AUC results is done with the ranking technique presented in [1]. We compute the mean value over five simulations for each algorithm. Next, we calculate ranks in every iteration of active learning for algorithms with the same batch size and datasets. The methods are compared to the aggregated mean ranks. The mean values are computed for ranks of 100, 200, ..., 1000 sampled instances. The reason of aggregation through a subset of all ranks is the understanding if more active learning iterations with smaller batch size or one with a larger one is better. The aggregated mean ranks for different datasets and batch sizes are in figure ...

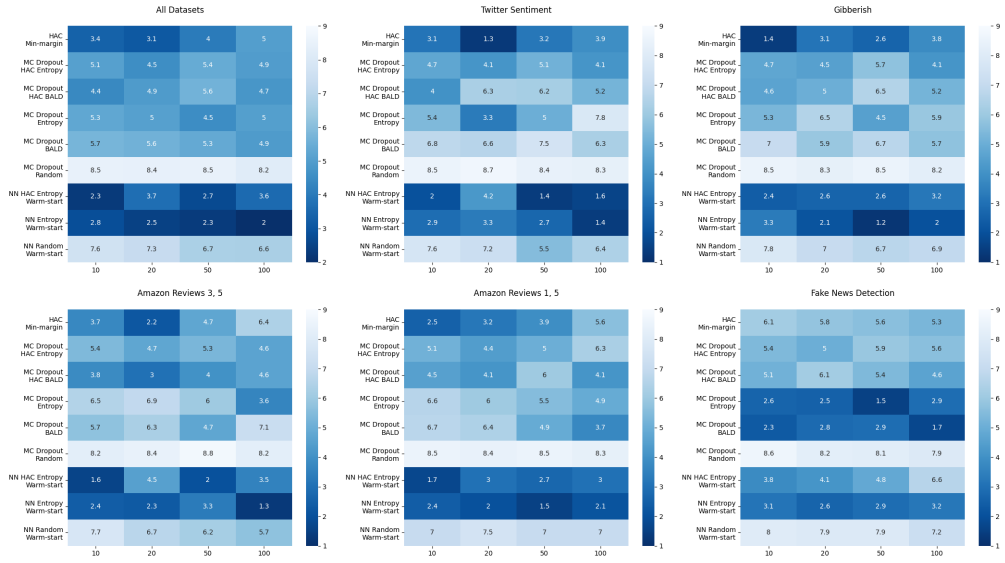


Figure 1

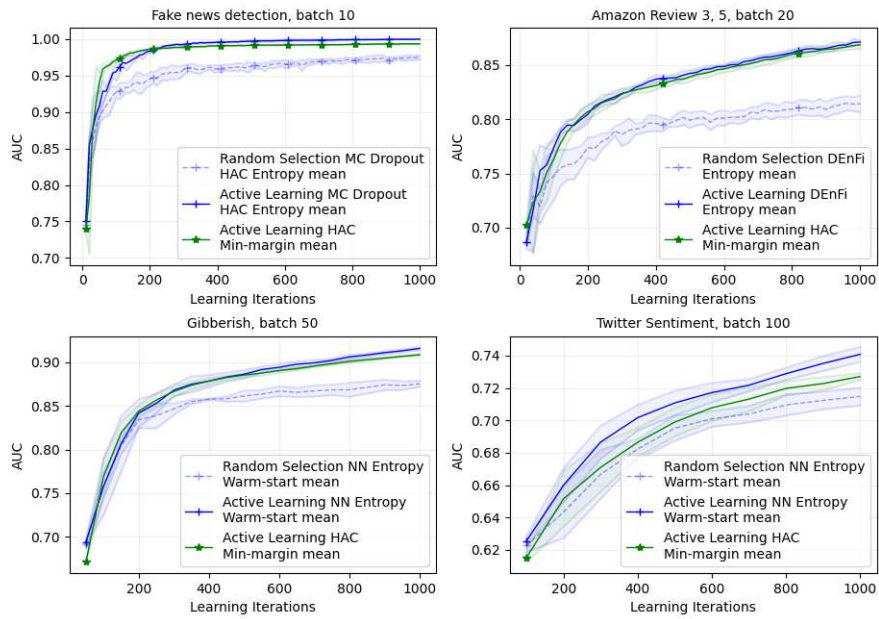


Figure 2